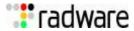
# Application and Network Orchestration Using HEAT & TOSCA

**Nati Shalom** 



Samuel Bercovici



@natishalom

@samuelbercovici



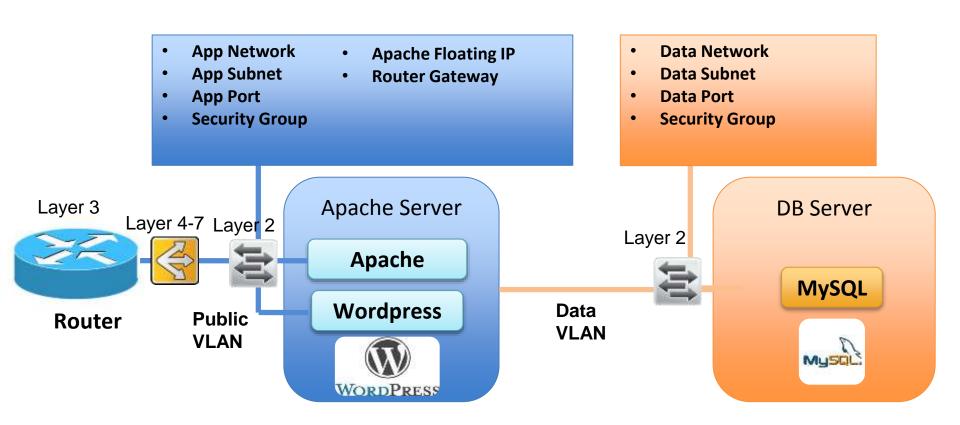
# Agenda

- Introduction
  - Introduction to Networking & Neutron
  - Introduction to Orchestration with HEAT & TOSCA
- Orchestrating Network & Applications by Example
  - Heat + Ceilometer Auto-Scaling Demo
  - Portable Version using TOSCA
- Real World (NFV) Example
  - Creating "Your own Skype" On-demand.
- Summary

# Let's Start With a Quick Overview...



# **Networking & Application**



#### What is Neutron?

Tenant Facing API to manage L2-L7 network logical elements

- Layer 2 networks
- Layer 3 subnets
  - IP address management DHCP based
  - Router / gateway / NAT
  - Port (Security groups, Floating IP)
- Layer 4-7 Services Load balancing, VPN, Firewall

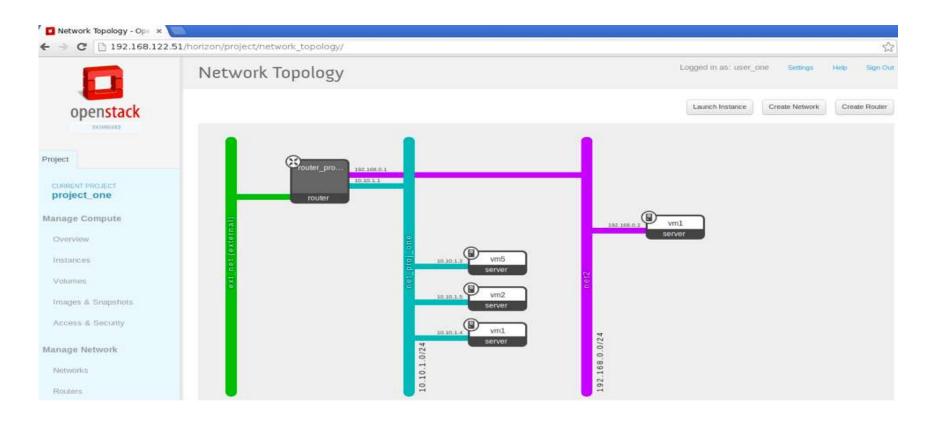




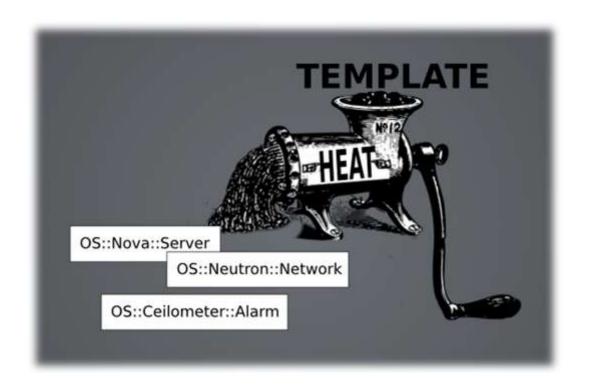




# **Networking in OpenStack Horizon**

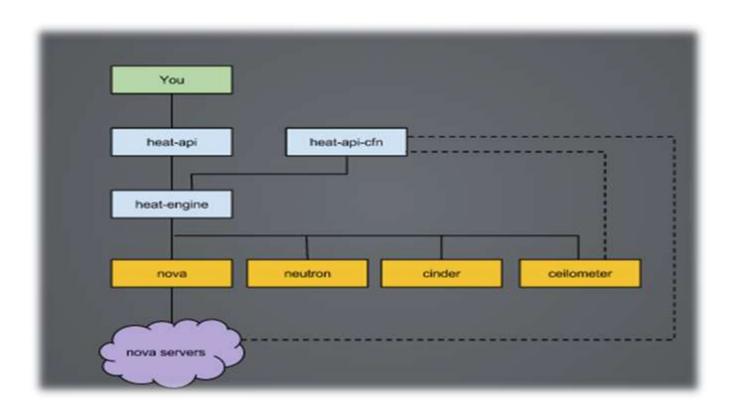


#### What is Heat?

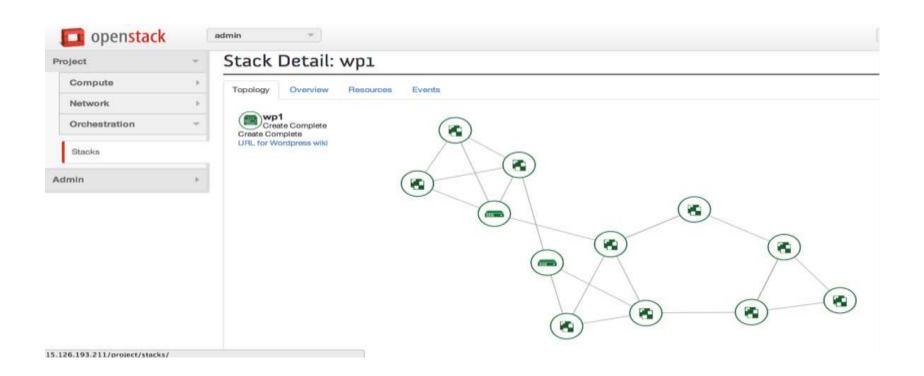


Heat provides a mechanism for orchestrating OpenStack resources through the use of modular templates.

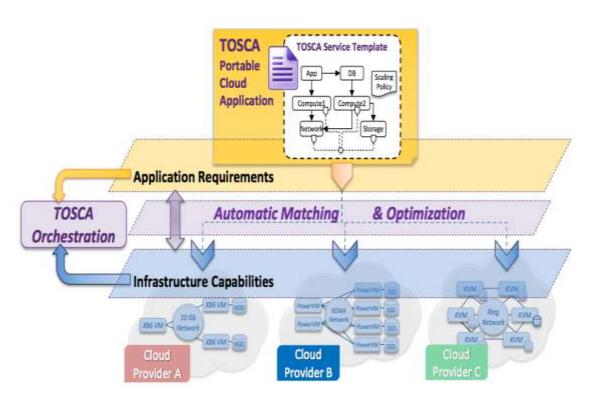
#### **Heat Architecture**



# **Heat Topology View**



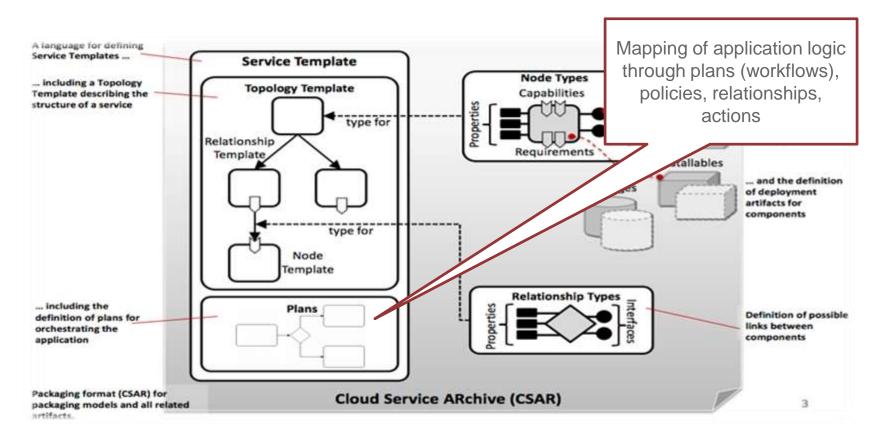
#### What is TOSCA?



TOSCA defines the interoperable description of applications; including their components, relationships, dependencies, requirements, and capabilities....



#### TOSCA in a Nutshell



#### **TOSCA State of the Union**

- Top four cloud open standard (Forrester)
- 5000+ participants
- 65+ countries
- TOSCA Parser Integrated into the Heat Project



The focus of this session



## **Comparing TOSCA & Heat**

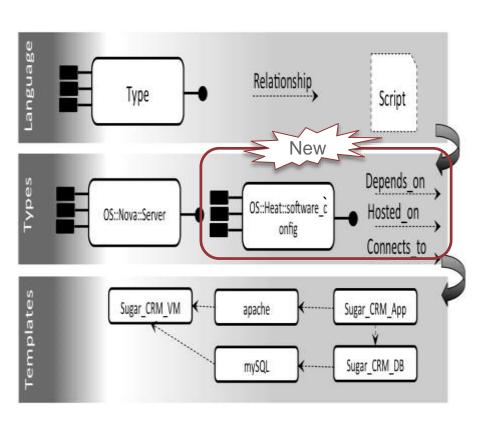
- Heat Automate the configuration and setup of OpenStack resources
- Specific to OpenStack
- TOSCA Automation of the application deployment and management lifecycle
- Portable

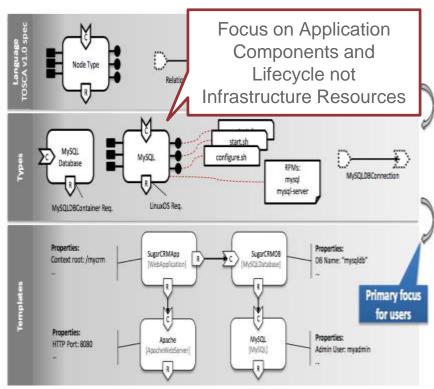


Merging Concepts



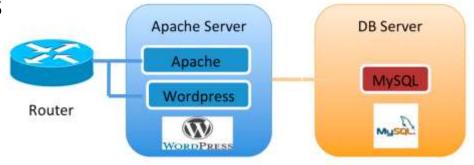
## **Heat/HOT vs TOSCA Syntax**





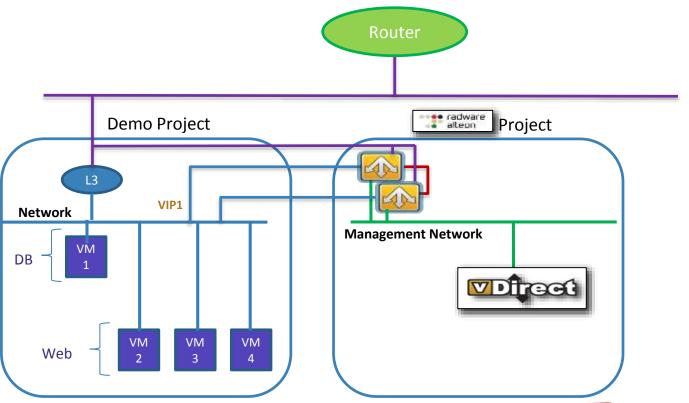
# Orchestrating Networks & Applications by Example...

- Define Networks as Parameters
- Create a New WordPress Stack
- WordPress Scale-out
- WordPress Scale-in



Reference Hot Templates: <u>autoscaling.yaml</u> and <u>lb\_server.yaml</u>

#### **Demo Scenario**





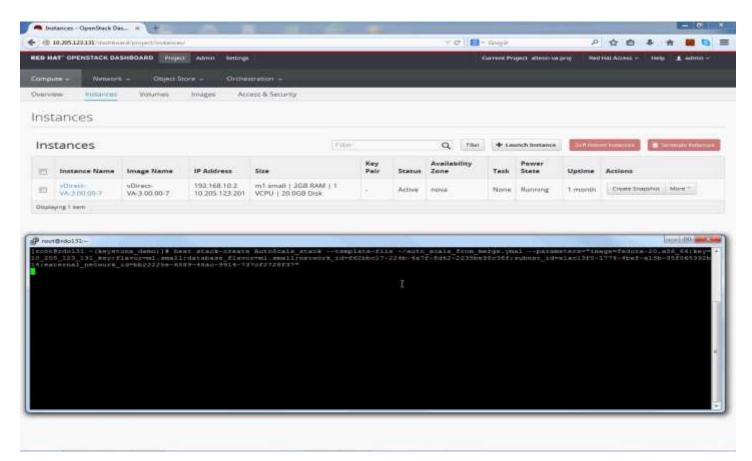
#### **Auto Scaling - Key Elements**

```
web server group:
  type: OS::Heat::AutoScalingGroup
 properties:
   min size: 1
   max size: 3
    resource:
      type: 1b server.yaml
     properties:
web server scaleup policy:
  type: OS::Heat::ScalingPolicy
  properties:
web server scaledown policy:
  type: OS::Heat::ScalingPolicy
  properties:
cpu alarm high:
  type: OS::Ceilometer::Alarm
  properties:
cpu alarm low:
  type: OS::Ceilometer::Alarm
  properties:
monitor:
  type: OS::Neutron::HealthMonitor
  properties:
pool:
  type: OS::Neutron::Pool
 properties:
1b:
  type: OS::Neutron::LoadBalancer
  properties:
# assign a floating ip address to the load balancer
# pool.
1b floating:
  type: "OS::Neutron::FloatingIP"
  properties:
```

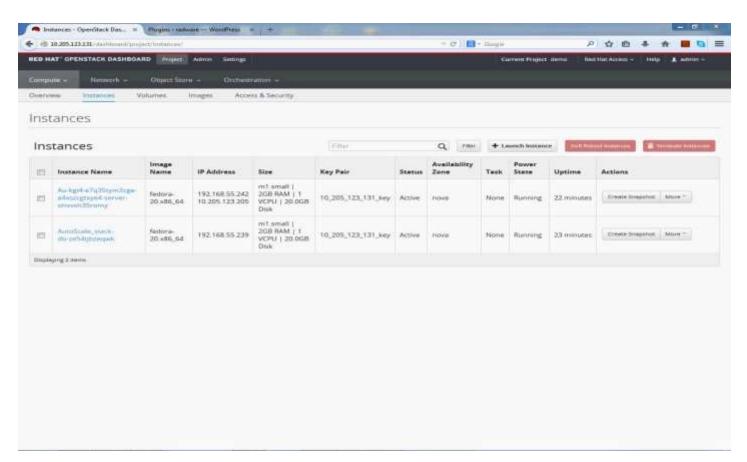
#### **Create New Stack**

```
#heat stack-create AutoScale stack
 --template-file ~/autoscale.yaml
 --parameters="image=fedora_image_for_web_and_db;
               key=some key;
               flavor=m1.small;database flavor=m1.small;
               network id=servers network id;
               subnet id=servers_subnet_id;
               external network=floating ip network
```

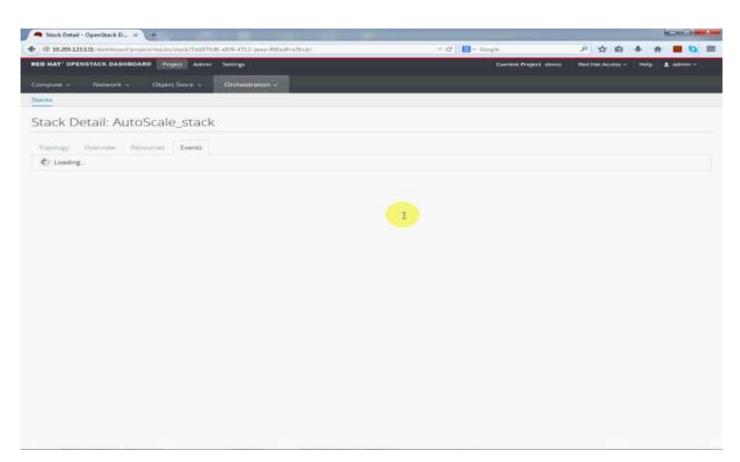
#### **Create New Stack**



#### **WordPress Scale-out**



#### WordPress Scale-In



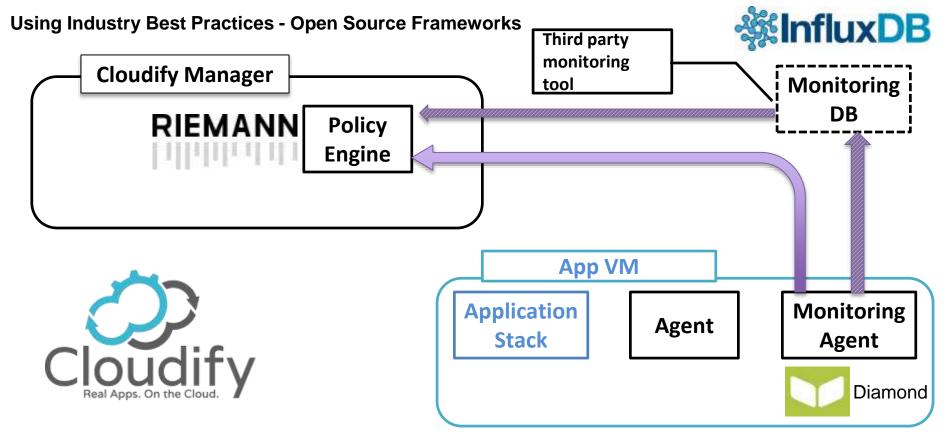
#### **Portable TOSCA Alternative**

```
node templates:
  wordpress:
    type: tosca.nodes.WebApplication.WordPress
    properties:
      admin user: { get input: wp admin username }
     admin password: { get input: wp admin password }
     db_host: { get_property: [ db_server, ip_address ] }
    requirements:
      - host: wordpress webserver
      - database_endpoint: wordpress_db
  wordpress webserver:
   type: tosca.nodes.WebServer.Apache
    properties:
        http port: 8080
        https_port: 8443
    requirements:
      - host: web server
    interaces:
     Lifecycle:
        inputs:
          db_host: { get_property: [ db_server, ip_address ] }
         db_port: { get_property: [ wordpress_db, db_port ] }
         db_name: { get_property: [ wordpress_db, db_name ] }
         db user: { get property: [ wordpress db, db user ] }
         db password: { get property: [ wordpress db, db password ] }
        configure: scripts/installandconfigurewp.sh
  web server:
    type: tosca.nodes.Compute
    properties:
      # Compute properties
      num cpus: 2
      mem size: 4
      disk size: 10
      # host image properties
      os_arch: x86_64
```

os\_type: linux os\_distribution: ubuntu os version: 12.04 Software components are defined just like any other resource

Infrastructure resource matched by their properties

# Portable Auto-Scaling



# Portable TOSCA (Like) Metrics

```
wordpress_vm
    type: cloudify.openstack.server
    properties:
       server: { get_input: server }
       image: { get_input: image }
       flavor_name: { get_input: flavor_name }
       security_groups: { get_input: security_groups }
    relationships:
          type: cloudify.relationships.depends_on
            target: Everything Allowed
    interfaces:
        cloudify.interfaces.monitor_lifecycle:
            - start:
                mapping: diamond.diamond_agent.tasks.install
                properties:
                    diamond_config:
                    interval: 5
                    collectors:
                        CPUCollector: {}

    stdp: diamond.diamond_agent.tasks.uninstall
```

**Set Diamond Collector** 



# Portable TOSCA (Like) Auto-Scaling

```
groups:
    wordpress group:
        members: [wordpress app]
        policies:
            monitoring_policy:
                type: threshold_policy
                properties:
                    metric: CPU PERCENTAGE
                    upper bound: 50
                    lower bound: 15
                triggers:
                    scaleup trigger:
                        type: cloudify.policies.triggers.execute workflow
                        parameters:
                            workflow: scaleup
                            workflow parameters:
                                 node_id: { get_properly: [SELF, node id ] }
                                 scale by: 1
                    scaledown trigger:
                        type: cloudify.policies.triggers.execute_workflow
                        parameters:
                            workflow: scaledown
                            workflow parameters:
                                 node_id: { get_properly: [SELF, node_id ] }
                                 scale by: 1
```

Set Auto-Scaling Policy
Based on the App
Metric



# **Mixing TOSCA & Heat**

**TOSCA Orchestration (Portable Lifecycle Orchestration)** 

Other Cloud Orchestration (Infrastructure Orchestration)

**Heat Orchestration (Infrastructure Orchestration)** 

**Other Cloud** 

**OpenStack** 

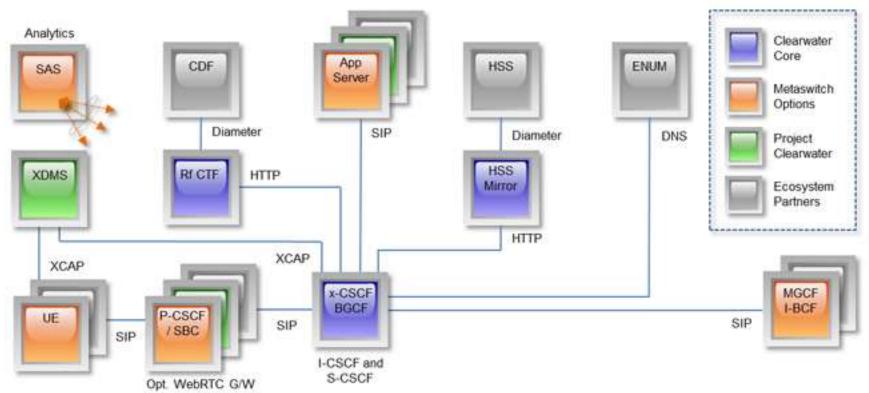
# Real Life (NFV) Use Case

Create Your Own "Skype" On-demand ...

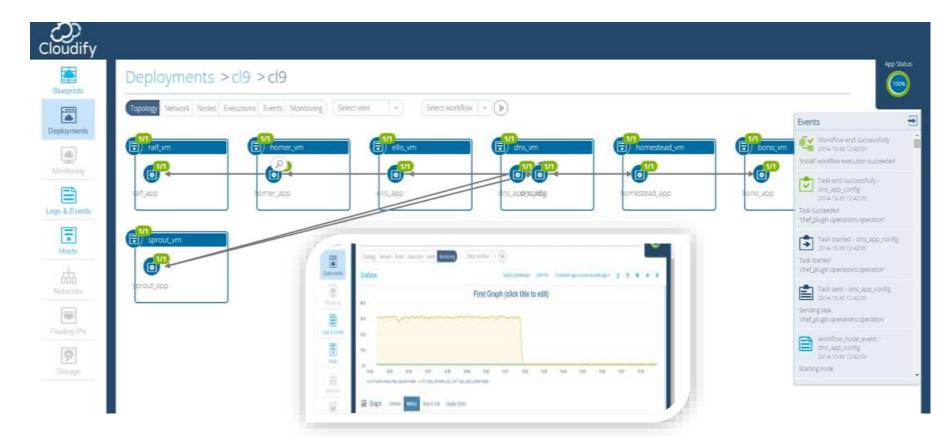








## **Automating the Entire Stack**



### **Summary**

Putting networks & apps together enables:

- Simplify the deployment of complex apps
- Tighten the security
- Increase the agility

Available Today: Heat & Neutron, Ceilometer TOSCA makes it portable

#### References

- Auto-Scaling WordPress Example (by Radware)
- ClearWater (NFV) Example (by Cloudify)
- Heat on DevStack
- Putting Heat & TOSCA Together
- Heat Technical Overview
- TOSCA Overview
- TOSCA Translator Project on StackForge