



**Svetlin Nakov**

**Manager Training  
and Inspiration**

**Software University**

**<http://softuni.bg>**

# HTTP/2 for Developers

## How HTTP/2 Works and How It Changes Developer's Life?



# http://2


# About Me

- Co-founder of the Software University (SoftUni) – <https://softuni.bg>
  - Trained thousands people in programming and software development
  - Taught hundreds of software development lessons
- 20+ years in software development
  - 15 years in training programmers
- Author of 7 programming books
  - Speaker at hundreds talks, presentations, lessons, courses, seminars, conferences, and technical events
- Blog: <http://www.nakov.com>



# Top 1M Alexa Sites (in 2013) – Telemetry

- Networking takes ~ 70% of the page load time!



- 69.5% of time blocked on network
- 6.6% of time blocked JavaScript
- 5.1% blocked on Layout
- 4.5% blocked on Paint

# HTTP/2: The Problem

- Opening a typical site in 2014, e.g. <https://softuni.bg>

65 requests | 269 KB transferred | Finish: 9.43 s | DOMContentLoaded: 2.45 s | Load: 2.78 s

- HTTP 1.0 supports only one connection per request
  - 65 connections need to be open, then closed
- HTTP 1.1 supports "keep alive"
  - Open connection, download a resource, then another resource, ...
  - Works in **serial mode**: if one resource is slow, page load is slow
- Modern Web needs "a better HTTP protocol"



# What is HTTP/2?


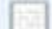




- HTTP/2 is protocol designed for
  - Low latency transport of content over the Web
- No change to HTTP semantics
  - It's about how data travels through the wire
- Key new features in HTTP/2
  - Multiplexing: multiple streams over a single connection
  - Header compression: reuse headers from previous requests
  - Server push: multiple parallel responses for a single request
  - Prioritization: some resources have priorities







# HTTP/2 History

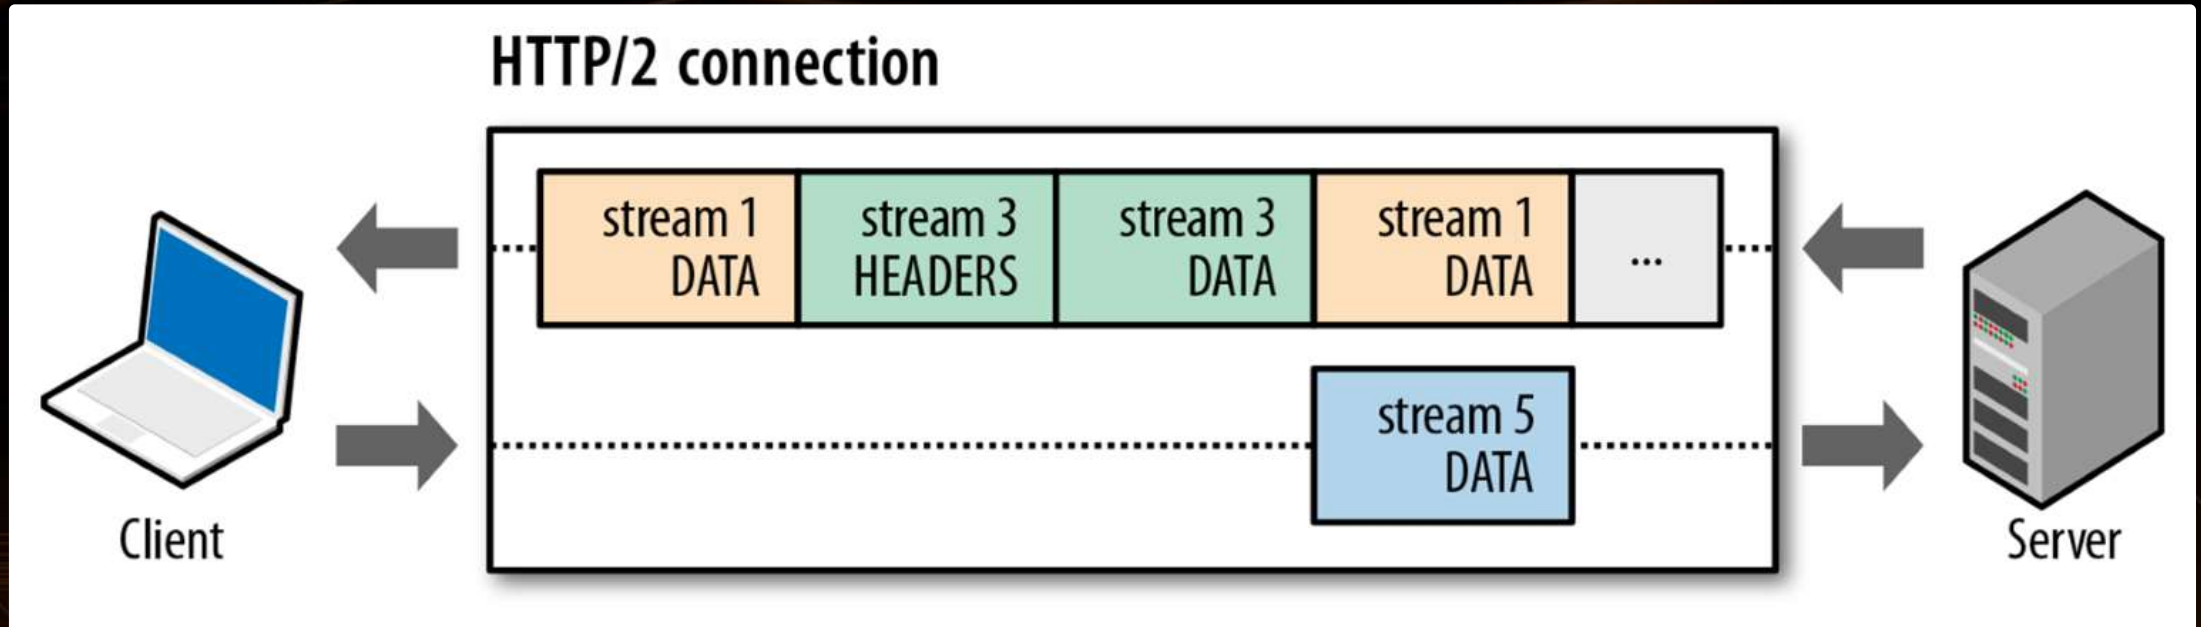
- **HTTP 0.9** (1991) – supports simple GET requests
- **HTTP 1.0** (1996) – single connection, POST, headers, ...
- **HTTP 1.1** (1999) – keep-alive connections + pipelining
- **SPDY** (2009-2010) – connection multiplexing + server push
  - Google developed the SPDY to replace HTTP 1.1 in Chrome and in the Google ecosystem
- **HTTP/2** (2014-2015) – de facto standardized SPDY
  - An official standard for the new low-latency HTTP protocol

# HTTP/2 Enabled Web Sites – Examples

- Facebook
  - Technically not HTTP/2
  - SPDY/3.1
- Google and its sites (GMail, YouTube, ...)
  - Use the latest HTTP/2
  - Combines QUIC + SPDY
- Twitter, Yahoo, Wikipedia

Name	Method	Status	Protocol	Scheme
 www.facebook.com	GET	307	http/1.1	http
 www.facebook.com	GET	200	h2-14	https
 mksngwPB2Xb.css	GET	200	spdy/3.1	https
 I5kTXq1bSJZ.css	GET	200	spdy/3.1	https
 hxxhYuhaB85C.css	GET	200	spdy/3.1	https
 MLi5OON7_Yz.css	GET	200	spdy/3.1	https

 gcosuc		200	h2	xhr
 gen_204?v=3&s=webhp&latyp=csi&imc=3&imn=3&imp=...		204	h2	text/ht...
 frame?sourceid=1&hl=en&origin=https%3A%2F%2Fwww....		200	h2	docu...
 rs=AGLTcCPfZQTHgHAJ6IAyALxG-i-tYiDctA		200	h2	styles...
 rs=AltRSTOKextu3QAZfV0s_SKuW3vYp-SZuA		200	h2	script
 spinner_32_794cfa28f324131c58c8934183d55d25.gif		200	h2	gif

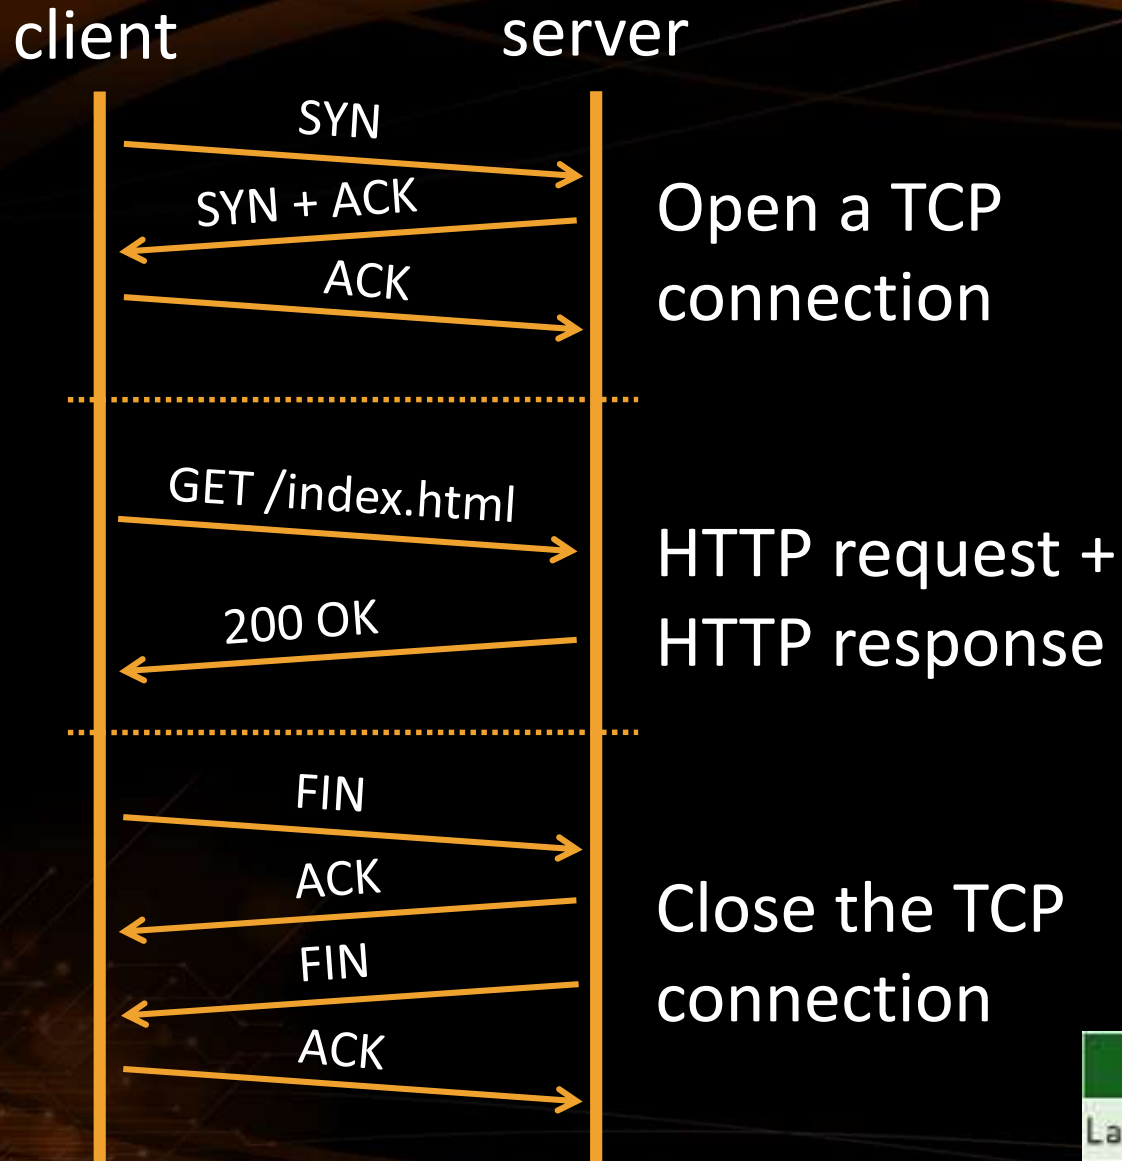


# Multiplexing

## Multiple Streams over Single Connection



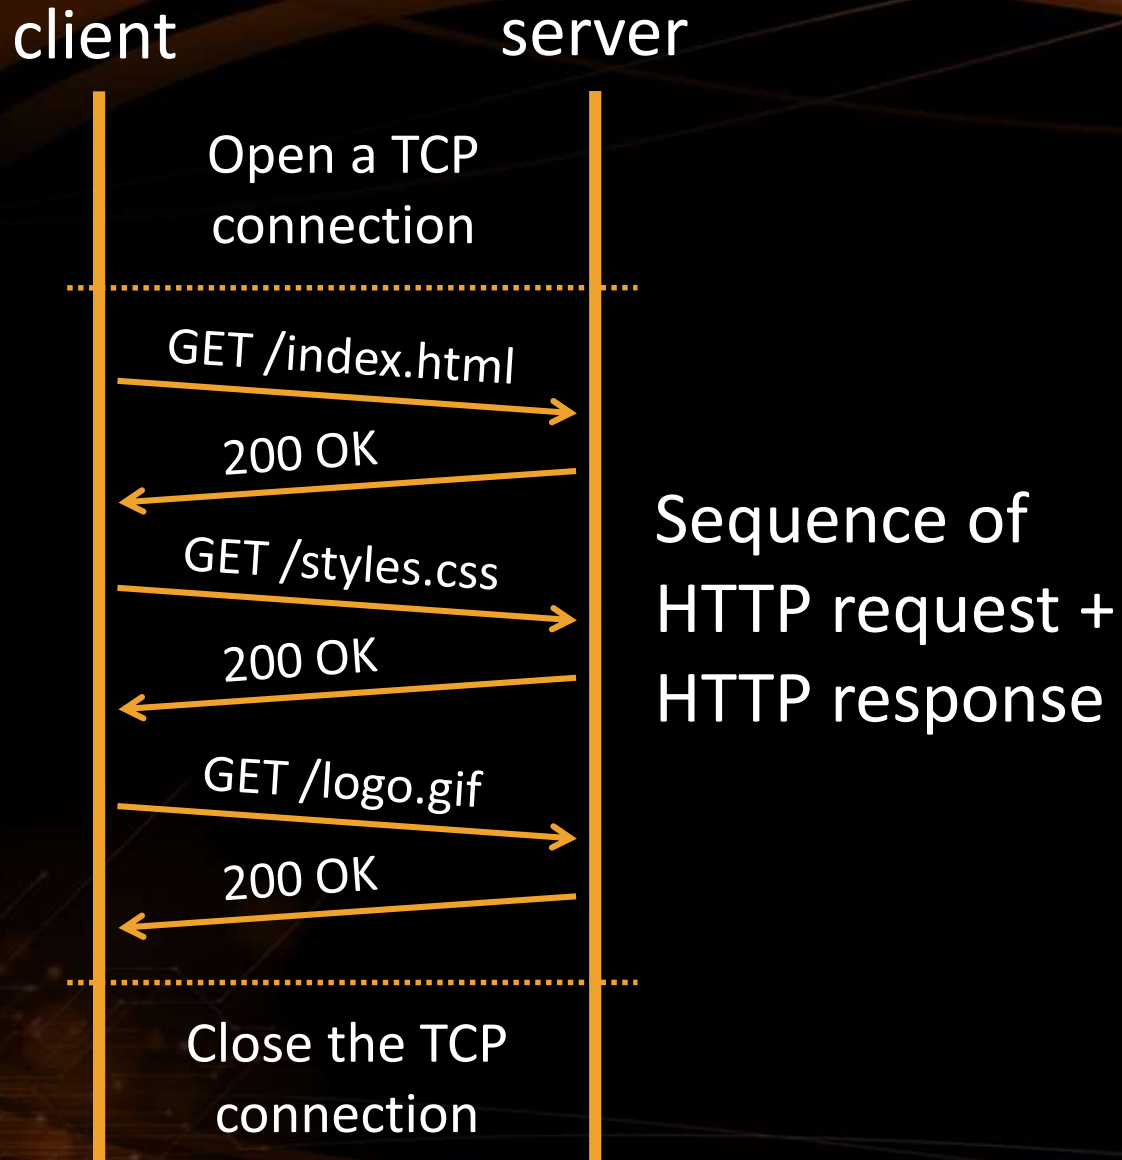
# HTTP 1.0: Single Connection



- HTTP 1.0 without **keep-alive**
  - Uses one request per TCP connection
  - Opening / closing a TCP connection is **slow**
    - Especially in high-latency connections (mobile networks)
- Typical latencies:

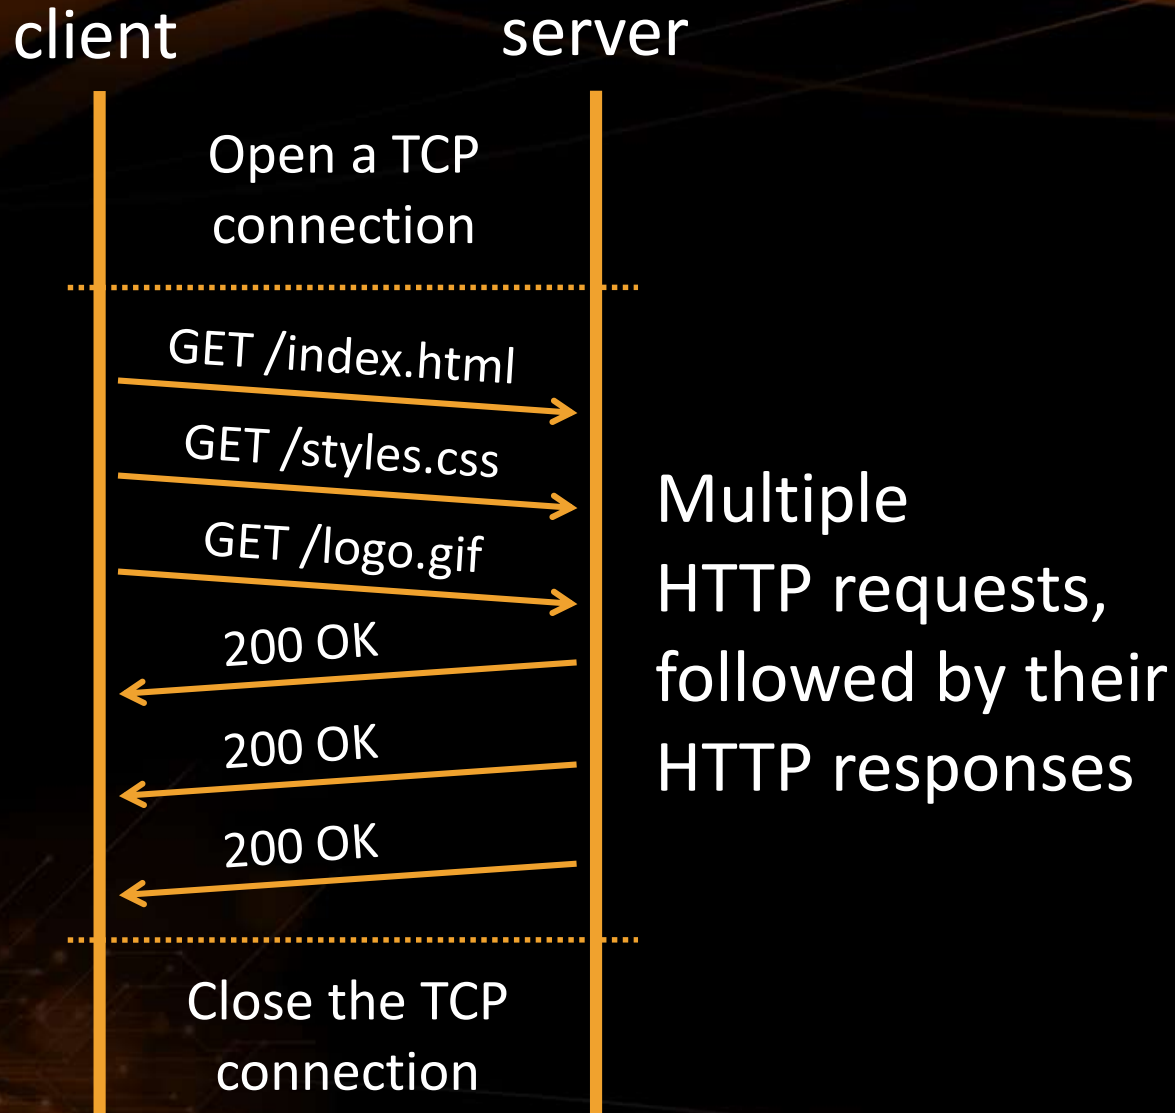
	LTE	HSPA+	HSPA	EDGE	GPRS
Latency	40-50 ms	100-200 ms	150-400 ms	600-750 ms	600-750 ms

# HTTP 1.1: Keep Alive



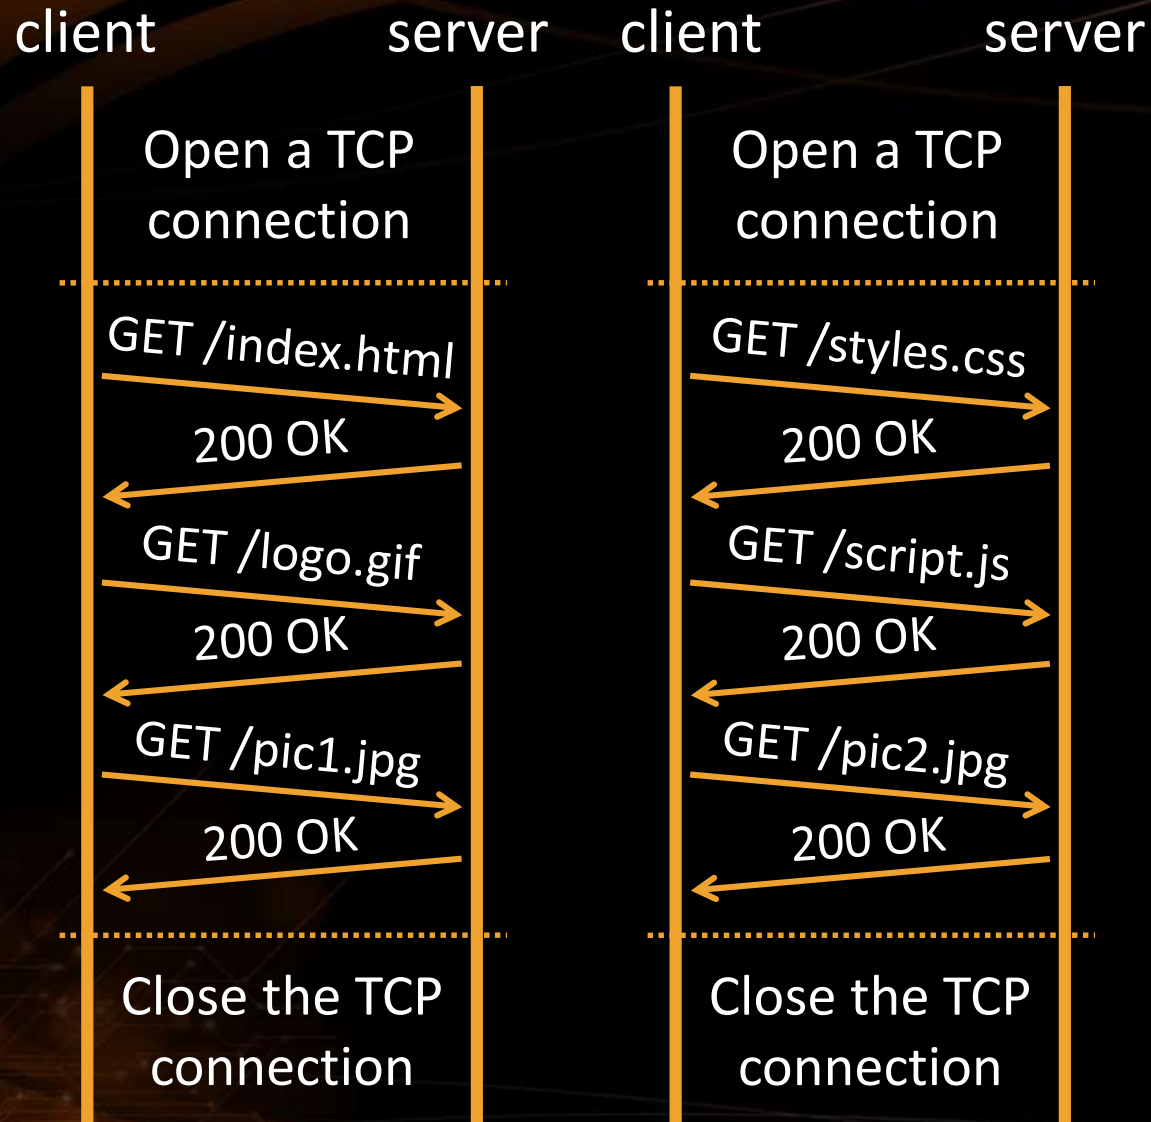
- HTTP 1.1 with keep-alive
  - Open a TCP connection
  - Multiple times perform:
    - Send a HTTP request
    - Read a HTTP response
  - Close the TCP connection
- Reuse TCP connections
  - Avoids multiple TCP connection open / close

# HTTP 1.1: Pipelining



- HTTP 1.1 with **pipelining**
  - Opens a TCP connection
  - Send multiple HTTP requests (without waiting responses)
  - Get the HTTP responses at once in FIFO order
  - Close the TCP connection
- **Head-of-line blocking (HOL)** slows down the entire pipeline

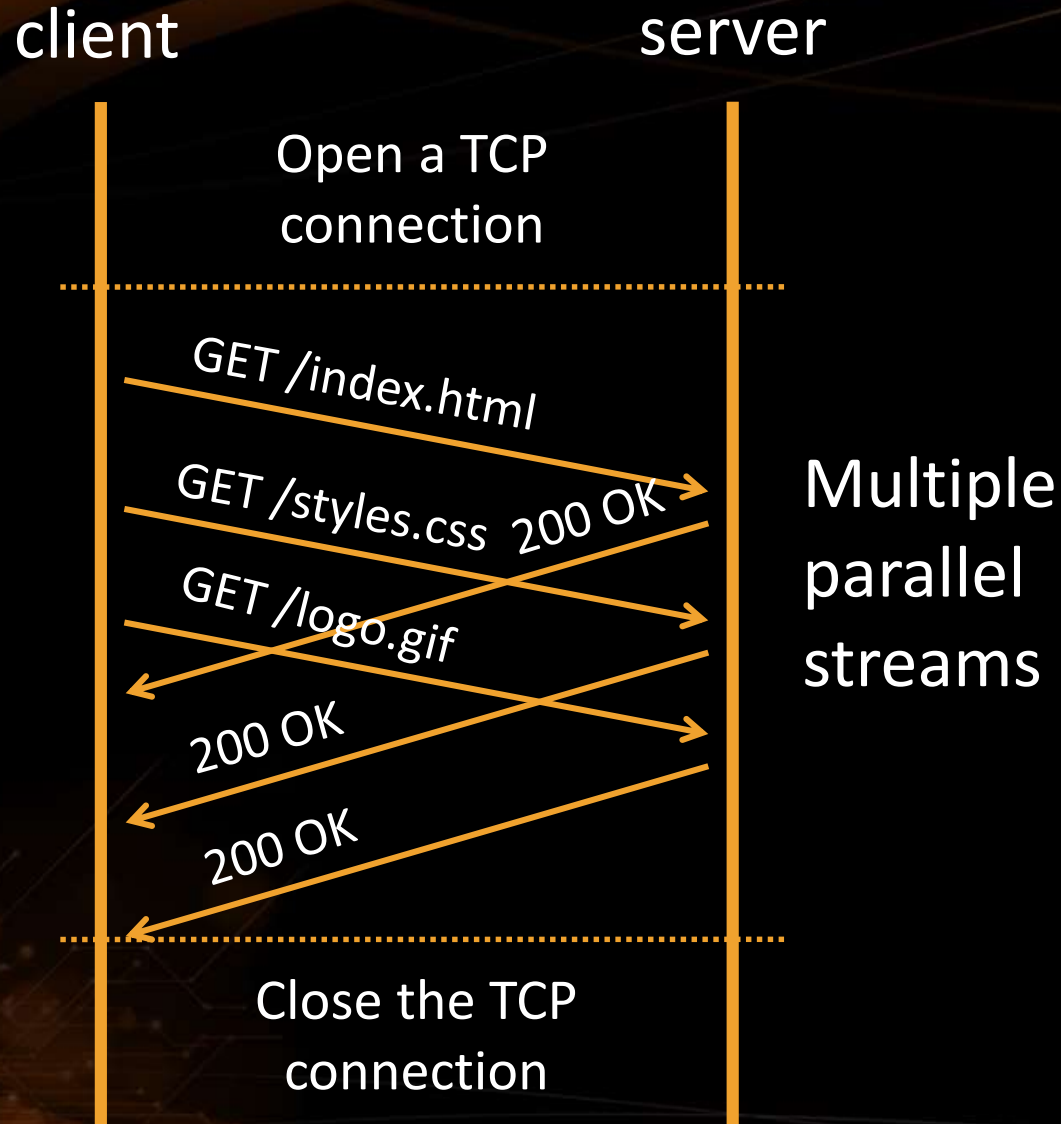
# HTTP 1.1: Multiple Connections



- Browsers open 4-8 parallel TCP connections to each server
  - Multiple request / response streams run in parallel
  - The initial opening of these connections still has a cost
- The number of connections is limited per domain
  - **Domain sharding** may improve load time for heavy sites




# HTTP/2: Multiplexing

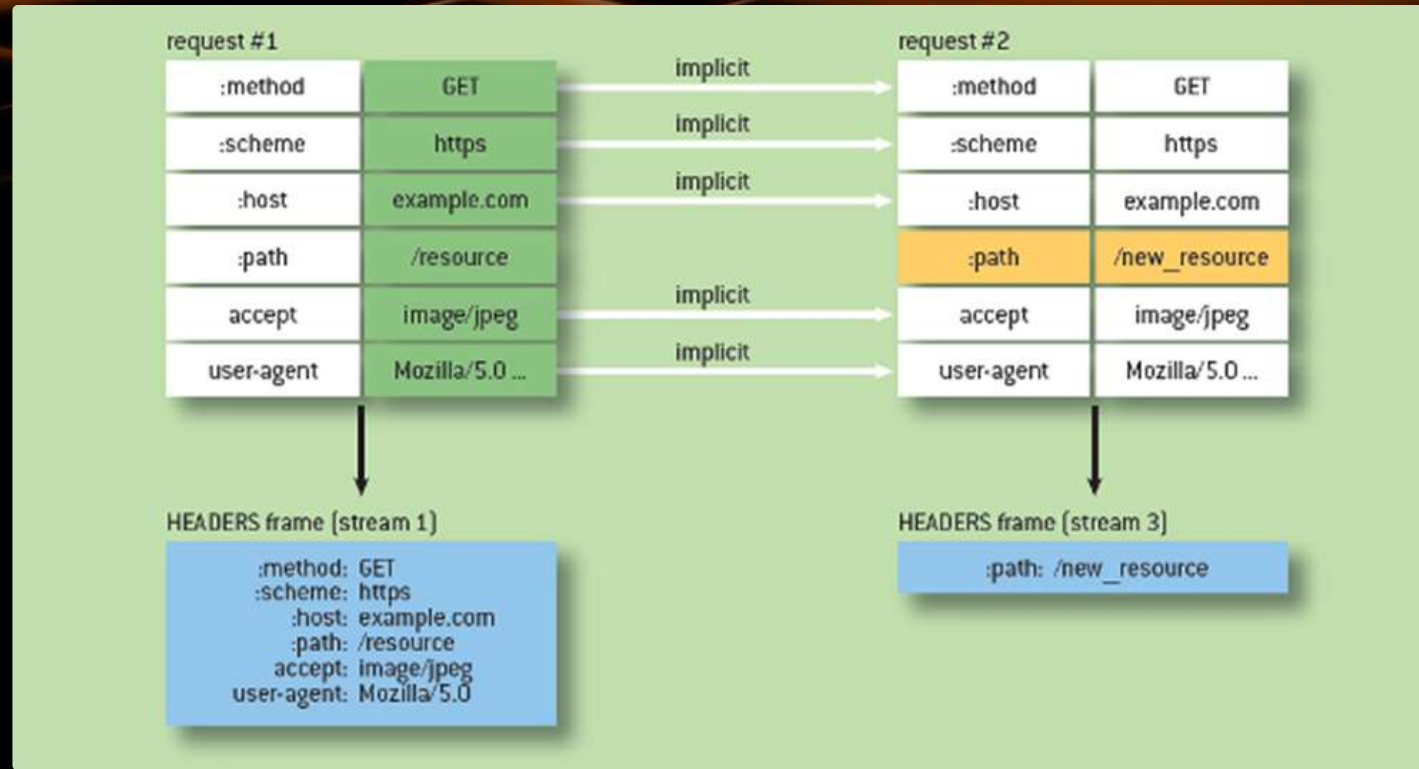


- HTTP/2 uses **multiplexing**
  - Single TCP connection
  - Multiple parallel streams
  - Prioritization: clients specify each resource importance
  - Out of order responses
- TCP connection is open once
- Resources come through the same connection

# How Multiplexing Works?

- HTTP/2 defines **streams** (bidirectional sequence of data)
  - One TCP connection serves multiple parallel streams
- The structure inside HTTP/2 is called a **frame**
  - Frame types: **HEADERS**, **DATA**, **SETTINGS**, **PUSH\_PROMISE**, ...





# HTTP Header Compression

## Reuse Headers from Previous Requests

# Typical HTTP Request Holds Long Headers

```
GET https://softuni.bg/ HTTP/1.1
Host: softuni.bg
Connection: keep-alive
Accept: text/html,application/xhtml+xml,application/xml,image/webp,*/*
Upgrade-Insecure-Requests: 1
User-Agent: Mozilla/5.0 (Windows NT 10.0; WOW64) AppleWebKit/537.36
(KHTML, like Gecko) Chrome/46.0.2490.86 Safari/537.36
Accept-Encoding: gzip, deflate, sdch
Accept-Language: en-US,en;q=0.8,bg;q=0.6
Cookie: .AspNet.SoftUniAuth=foQI2JhdsV0jE5PKhLAf1cM1B9b1_rN3Rmo3hdR...;
_ga=GA1.2.925942971.1440514837
```

**992 bytes**



# Next Request Holds Nearly the Same Headers

```
GET https://softuni.bg/apply HTTP/1.1
Host: softuni.bg
Connection: keep-alive
Accept: text/html,application/xhtml+xml,application/xml,image/webp,*/*
Upgrade-Insecure-Requests: 1
User-Agent: Mozilla/5.0 (Windows NT 10.0; WOW64) AppleWebKit/537.36
(KHTML, like Gecko) Chrome/46.0.2490.86 Safari/537.36
Referer: https://softuni.bg/
Accept-Encoding: gzip, deflate, sdch
Accept-Language: en-US,en;q=0.8,bg;q=0.6
Cookie: __RequestVerificationToken=5Eo-1AV9...1; language=bg;
_ga=GA1.2.925942971.1440514837; _gat=1;
.AspNet.SoftUniAuth=foQI2JhdsV0jE5PKhLAf1cM1B9b1_rN3Rmo3hdR...
```

**1173 bytes (181 bytes new)**

# Next Request Holds Even More Similar Headers

```
GET https://softuni.bg/contacts HTTP/1.1
Host: softuni.bg
Connection: keep-alive
Accept: text/html,application/xhtml+xml,application/xml,image/webp,*/*
Upgrade-Insecure-Requests: 1
User-Agent: Mozilla/5.0 (Windows NT 10.0; WOW64) AppleWebKit/537.36
(KHTML, like Gecko) Chrome/46.0.2490.86 Safari/537.36
Referer: https://softuni.bg/
Accept-Encoding: gzip, deflate, sdch
Accept-Language: en-US,en;q=0.8,bg;q=0.6
Cookie: __RequestVerificationToken=5Eo-1AV9...1; language=bg;
_ga=GA1.2.925942971.1440514837; _gat=1;
.AspNet.SoftUniAuth=foQI2JhdsV0jE5PKhLAflcMlB9b1_rN3Rmo3hdR...
```

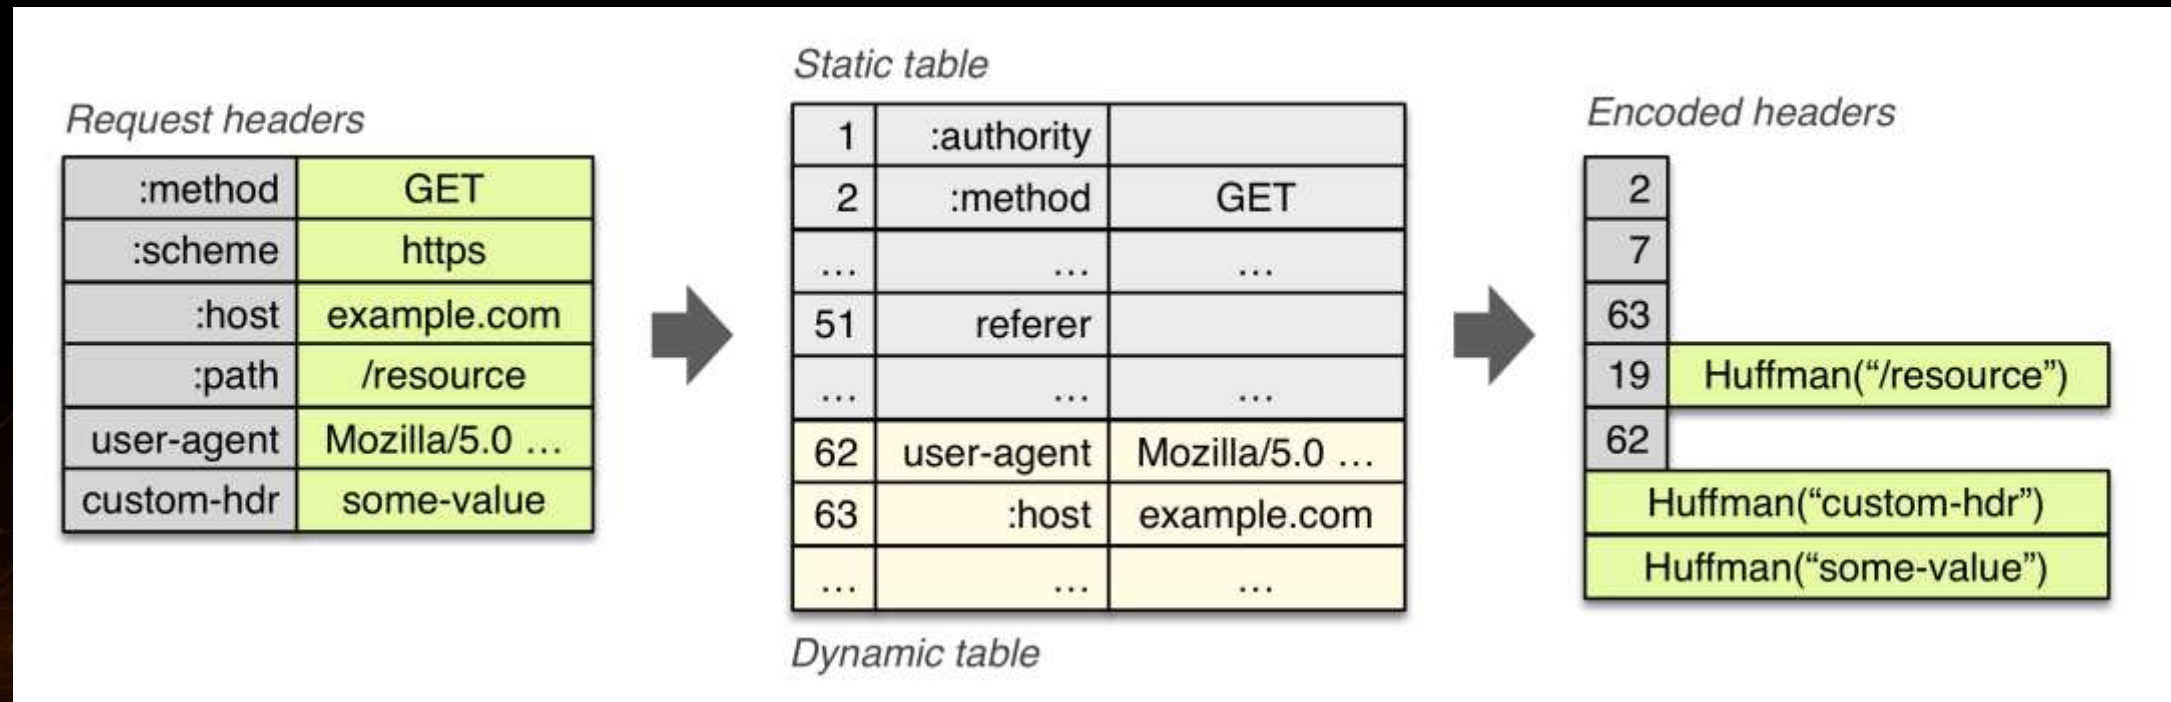
**1176 bytes (8 bytes new, headers 100% the same)**

# HPACK for HTTP/2 Headers Compression

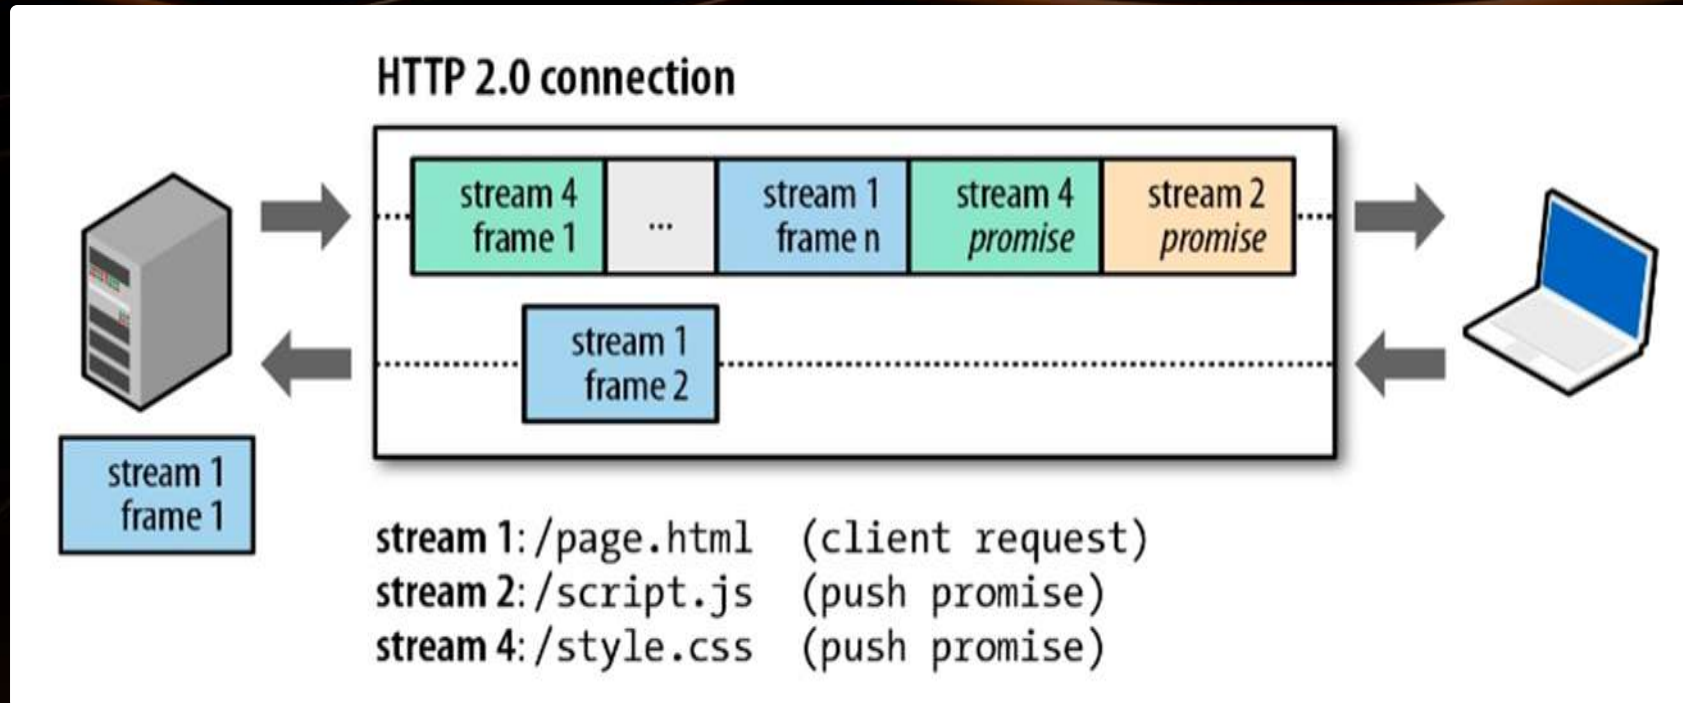
- HTTP/2 uses a compression standard called **HPACK** (RFC 7541)
  - Compresses the request and response HTTP headers
    - No overhead if headers do not change (e.g. polling requests)
  - Client and server maintain a **table of all previous headers**
    - Used to efficiently encode previously transmitted values
    - Header tables persist for the entire HTTP 2.0 connection
    - Incrementally updated by both the client and the server
  - Static (predefined) **Huffman code** used to compress values

# HPACK for HTTP/2 Headers Compression

- Request line is split to **pseudo headers** like **:method**, **:path**, ...
- **Static table** defines common HTTP headers
- **Dynamic table** defines HTTP headers in the current HTTP session





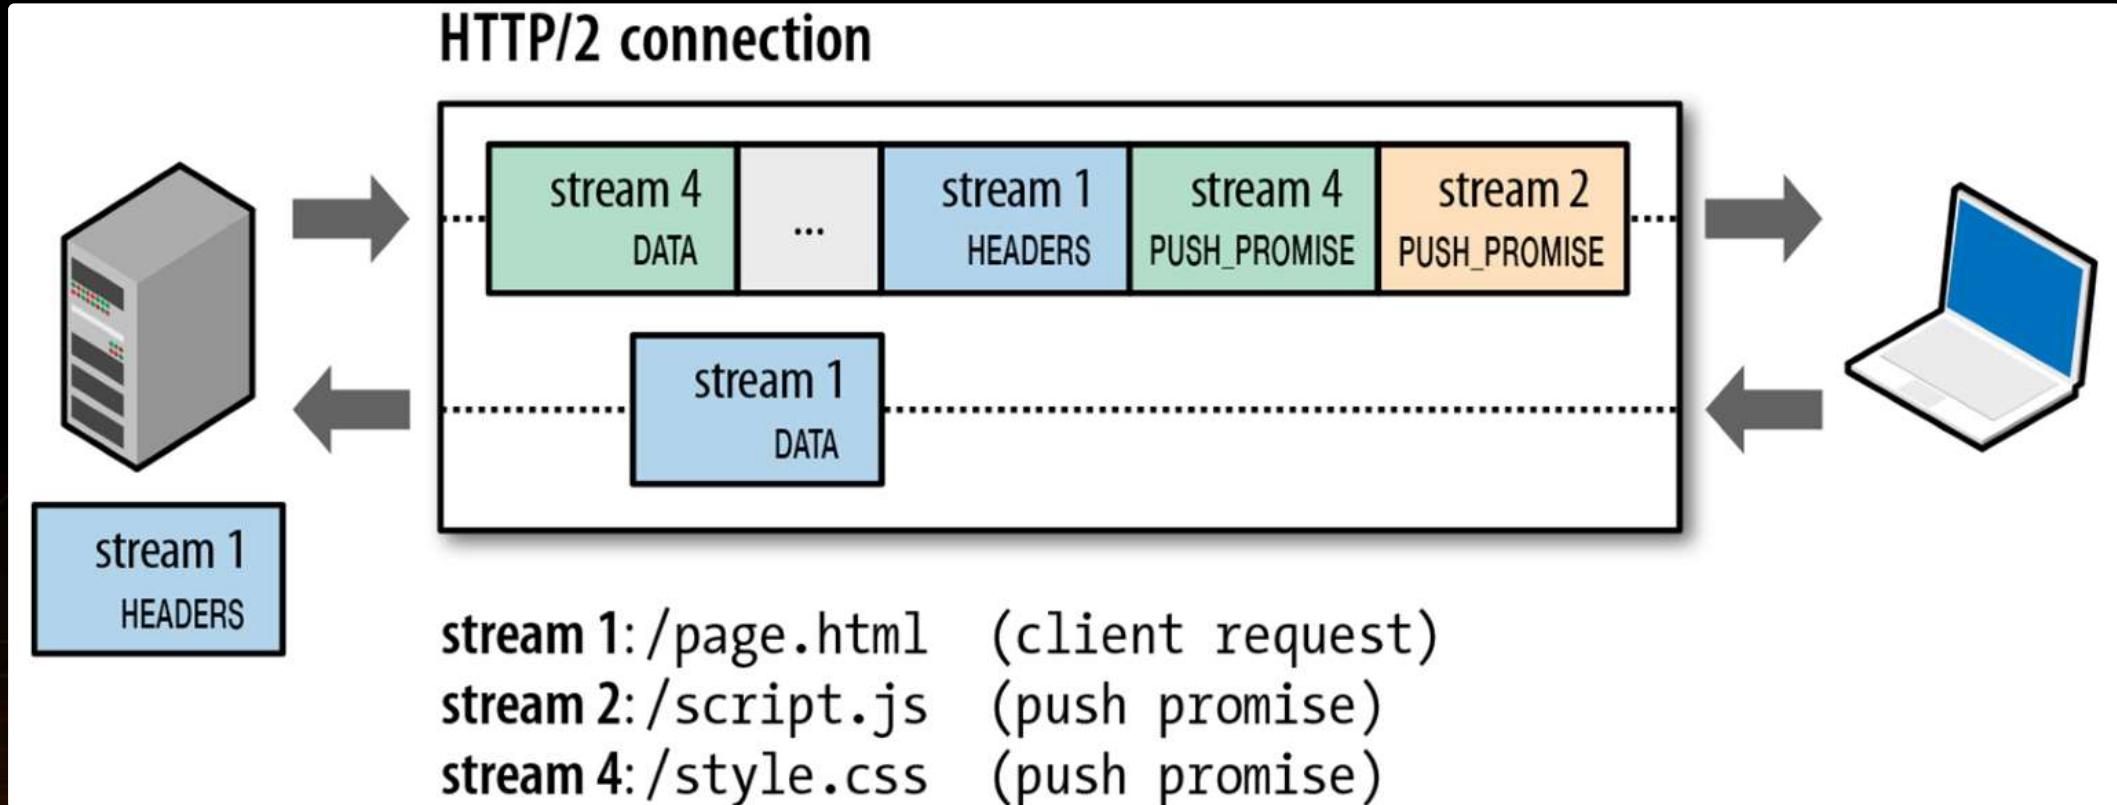


# HTTP/2 Server Push

## Multiple Parallel Responses for a Single Request

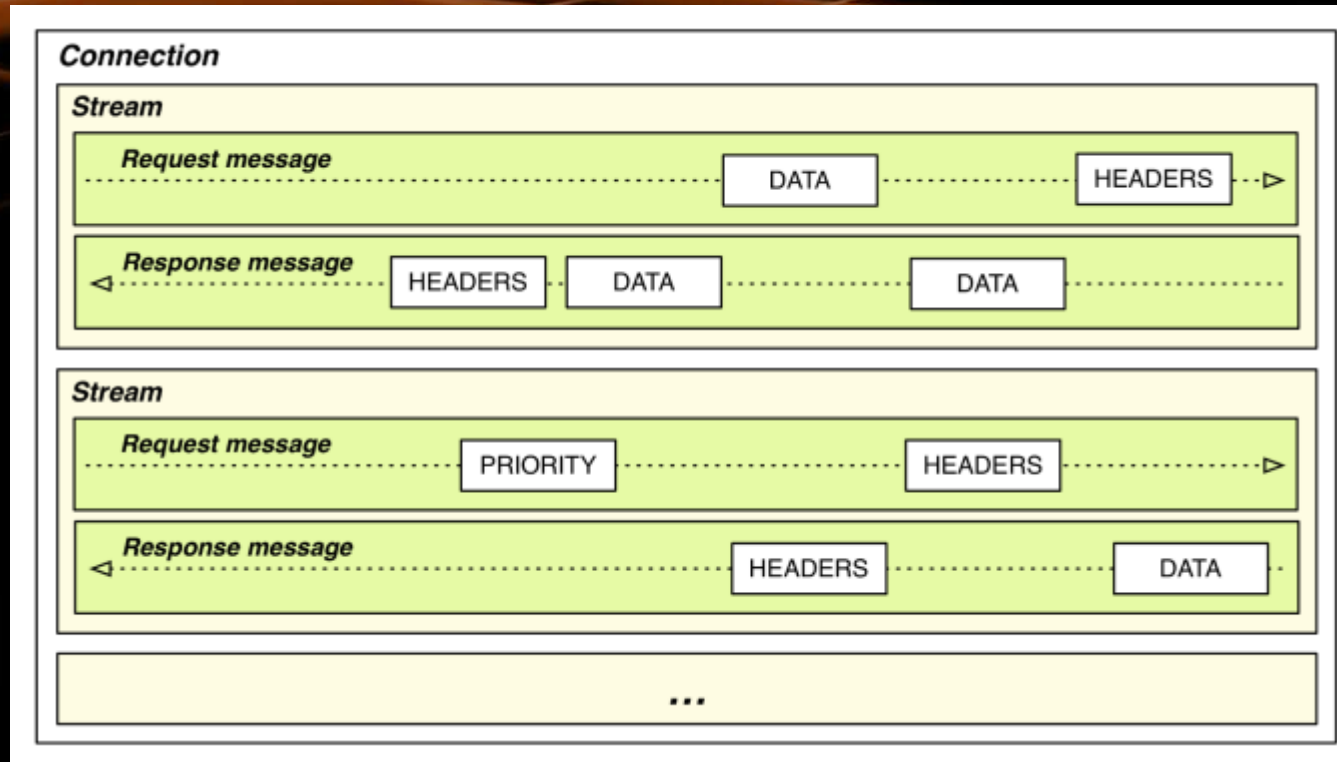
# Server Push in HTTP/2

- HTTP/2 servers can send **multiple responses** for a single client request (push additional resources with the requested resource)



# Push Promises

- **PUSH\_PROMISE** frames in HTTP/2 allows the server to send additional resources
  - First **headers** are sent, so **clients can decide** what to do
  - Clients can decline pushed resources (via a **RST\_STREAM** frame)
    - E.g. when the pushed resource is already in the cache
- Pushed resources can be cached and reused later
- Pushed resources can be multiplexed alongside other resources
- Pushed resources can be prioritized by the server



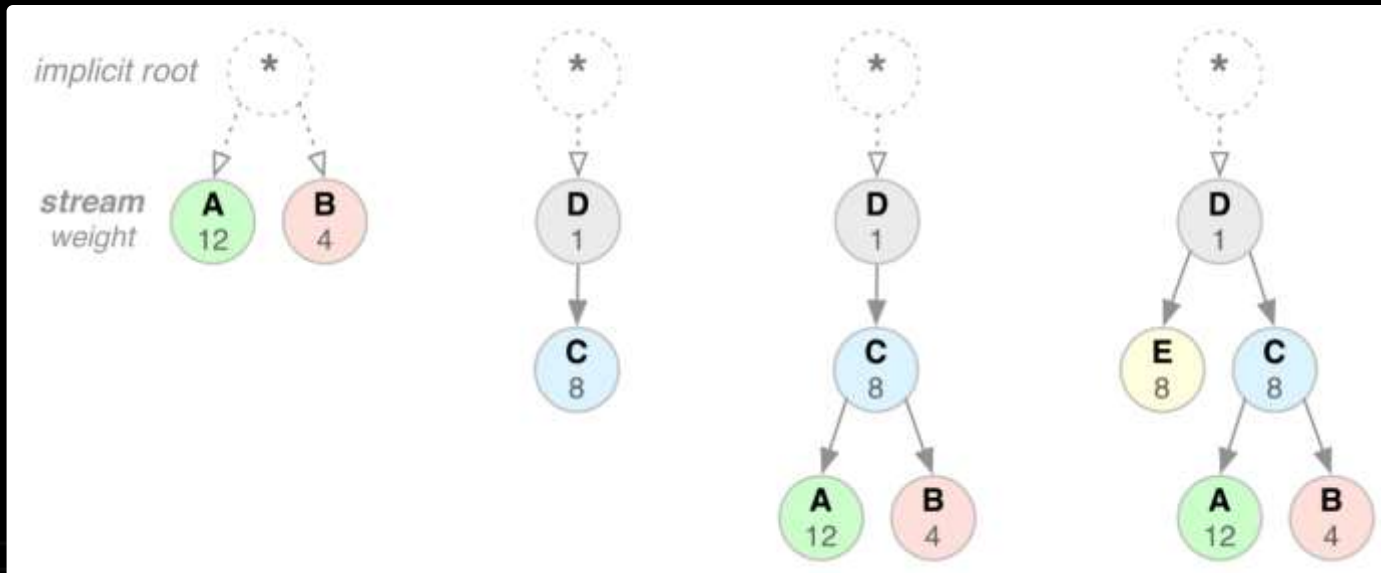
# Prioritization and Flow Control

## Weights, Dependencies, Flow Control



# Stream Weight and Dependency

- HTTP/2 defines stream prioritization by weight + dependency
  - Each stream may be assigned a **weight** between 1 and 256
  - Each stream may be given a **dependency** on another stream
- The "**prioritization tree**" defines the preferred order of resources

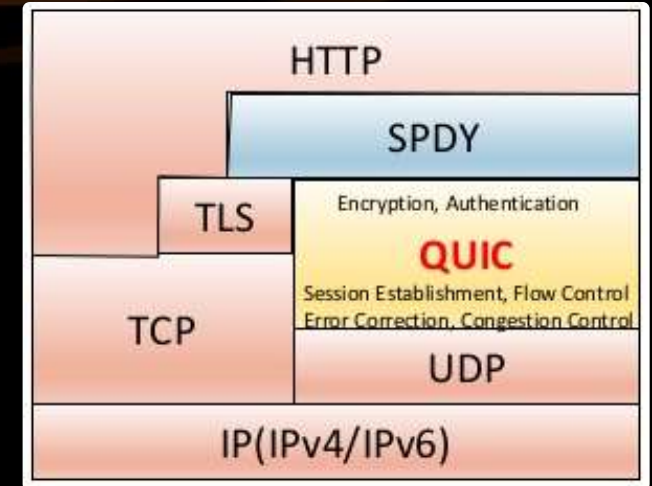


# Flow Control in HTTP/2

- What is **flow control**?
  - Managing the **rate of data transmission** between two nodes to prevent a fast sender from overwhelming a slow receiver
- Example:
  - User watches a video and presses the "pause" button
- For each HTTP/2 stream and for the entire connection
  - Receivers can specify **how many DATA bytes** is ready to receive
- **WINDOW\_UPDATE** frames configure the window size

# The QUIC Protocol

- Some browsers use the **QUIC** protocol
  - QUIC == Quick UDP Internet Connections
  - Low-latency UDP-based transport layer protocol
  - Designed to **replace TCP** for better performance
  - Build-in **multiplexing** and TLS-equivalent **security**
  - Still **experimental**, supported in Chromium
- Learn more at:
  - <https://www.chromium.org/quic>
  - <https://en.wikipedia.org/wiki/QUIC>



**QUIC**

means

Quick UDP (User Datagram Protocol) Internet Connections



# Using HTTP/2

How to Switch to HTTP/2?



# How to Establish HTTP/2 Connection?

- There is **no** such thing as:
  - <http2://mysite.com>
- HTTP/2 uses a standard <https://> connection
  - With **TLS ALPN** (Application Layer Protocol Negotiation)
  - **ALPN** allows the applications to negotiate which protocol should be used over a secure connection without round trips
  - Most browsers do not support HTTP/2 without SSL
- Browsers **check for HTTP/2 support** during the **SSL handshake**

# HTTP/2 Connection Upgrade

- HTTP 1.1 connection could be upgraded to HTTP/2:

```
GET / HTTP/1.1
Host: server.example.com
Connection: Upgrade, HTTP2-Settings
Upgrade: h2c
HTTP2-Settings: <base64url-encoded HTTP/2 SETTINGS>
```

- TLS is not mandatory for HTTP/2 (as defined in RFC 7540)
  - But most Web browsers refuse to implement clear text HTTP/2
- Non-browser client apps can still use clear-text HTTP/2



# HTTP/2 Tools

How It Changes Developer's Life?

# HTTP/2 is a Binary Protocol

- HTTP 1.1 was simple to test / debug / trace
  - It is **text-based** protocol, human-readable
- HTTP/2 will require tools
  - It is **binary protocol**, not human-readable
  - It runs over TLS tunnel, so sniffing is complicated

```
Command Prompt - ssh root@svn.softuni.org
root@svn:~# telnet softuni.org 80
Trying 98.124.199.71...
Connected to softuni.org.
Escape character is '^]'.
GET / HTTP/1.1
Host: softuni.org

HTTP/1.1 302 Found
Date: Sat, 21 Nov 2014 12:00:00 GMT
Content-Type: text/html
```



**h2i**


```
$ h2i google.com
Connecting to google.com:443 ...
Connected to 74.125.224.41:443
Negotiated protocol "h2-14"
[FrameHeader SETTINGS len=18]
  [MAX_CONCURRENT_STREAMS = 100]
  [INITIAL_WINDOW_SIZE = 1048576]
  [MAX_FRAME_SIZE = 16384]
[FrameHeader WINDOW_UPDATE len=4]
  Window-Increment = 983041
```



# Tools Not Working with HTTP/2

- **Wireshark** can sniff HTTP/2
  - But cannot decrypt TLS traffic
    - No out-of-the-box man-in-the-middle proxy
    - Needs the SSL private key for the remote site
  - In fact Wireshark is not of help for HTTP/2 Web developers
- **Fiddler** does not support HTTP/2
  - Wait for Microsoft to implement TLS ALPN in .NET Framework
  - It decrypts HTTPS through elegant **man-in-the-middle proxy**

# HTTP/2 and SPDY Indicator for Chrome



## HTTP/2 and SPDY indicator


offered by rauchg

★★★★★ (135) | [Developer Tools](#) | 55,675 users

OVERVIEW | REVIEWS | SUPPORT | RELATED

ADDED TO CHROME

G+1 420



Compatible with your device

An indicator in the address bar for HTTP/2 and SPDY support by each website.

HTTP/2 (based on Google's SPDY) makes the information exchange between browser and server significantly more performant. Websites and applications that upgrade their infrastructure to support them are at a clear advantage.

This extension shows a lightning icon in the address bar that's blue if the page is HTTP/2 enabled, green if the page is SPDY enabled, and optionally gray if neither is available.

- Indicates HTTP/2, SPDY, QUIC, versions, etc.

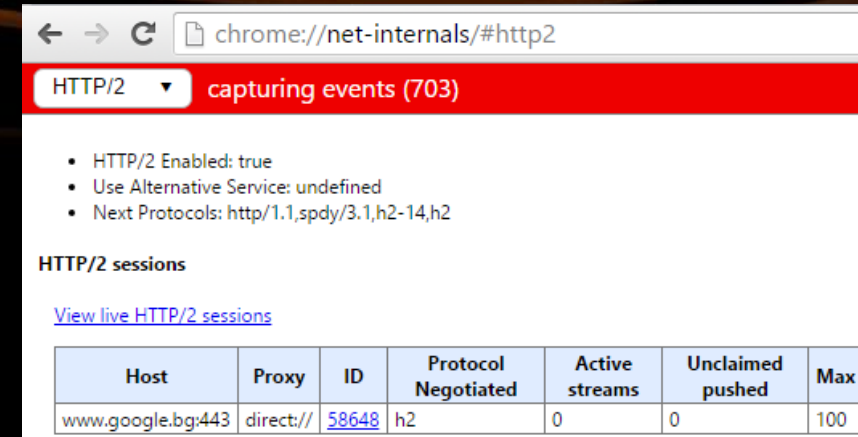
# Chrome HTTP/2 Tools

- Chrome HTTP/2 info page

- <chrome://net-internals/#http2>

- View all HTTP/2 sessions

- [chrome://net-internals/#events&q=type:HTTP2\\_SESSION](chrome://net-internals/#events&q=type:HTTP2_SESSION)



chrome://net-internals/#http2

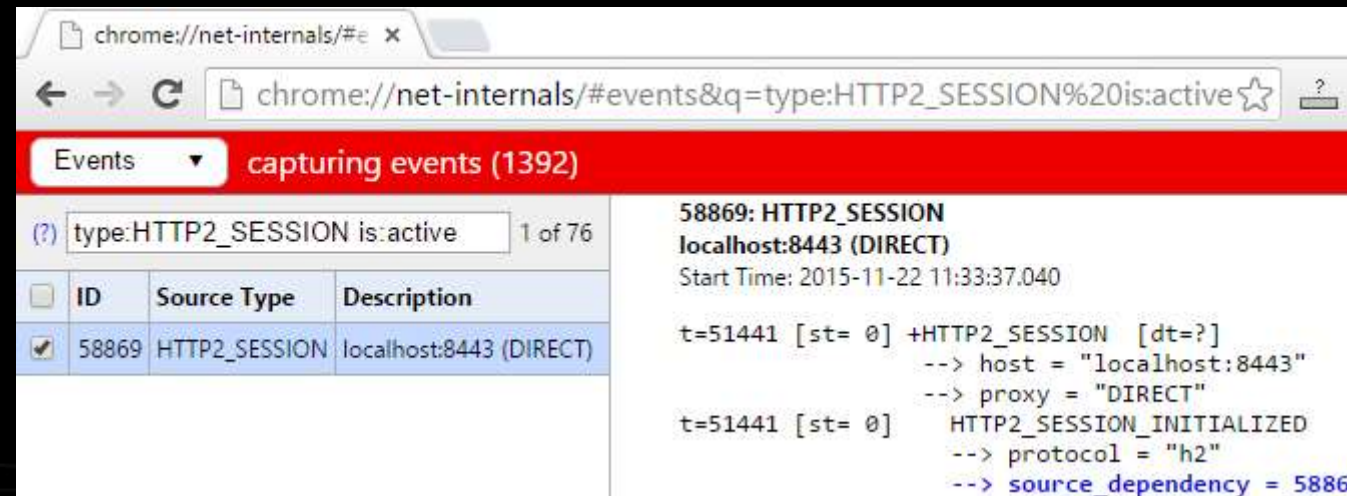
HTTP/2 capturing events (703)

- HTTP/2 Enabled: true
- Use Alternative Service: undefined
- Next Protocols: http/1.1,spdy/3.1,h2-14,h2

HTTP/2 sessions

[View live HTTP/2 sessions](#)

Host	Proxy	ID	Protocol Negotiated	Active streams	Unclaimed pushed	Max
www.google.bg:443	direct://	<a href="#">58648</a>	h2	0	0	100



chrome://net-internals/#e x

chrome://net-internals/#events&q=type:HTTP2\_SESSION%20is:active

Events capturing events (1392)

(?) type:HTTP2\_SESSION is:active 1 of 76

ID	Source Type	Description
<input checked="" type="checkbox"/> 58869	HTTP2_SESSION	localhost:8443 (DIRECT)

58869: HTTP2\_SESSION  
localhost:8443 (DIRECT)  
Start Time: 2015-11-22 11:33:37.040

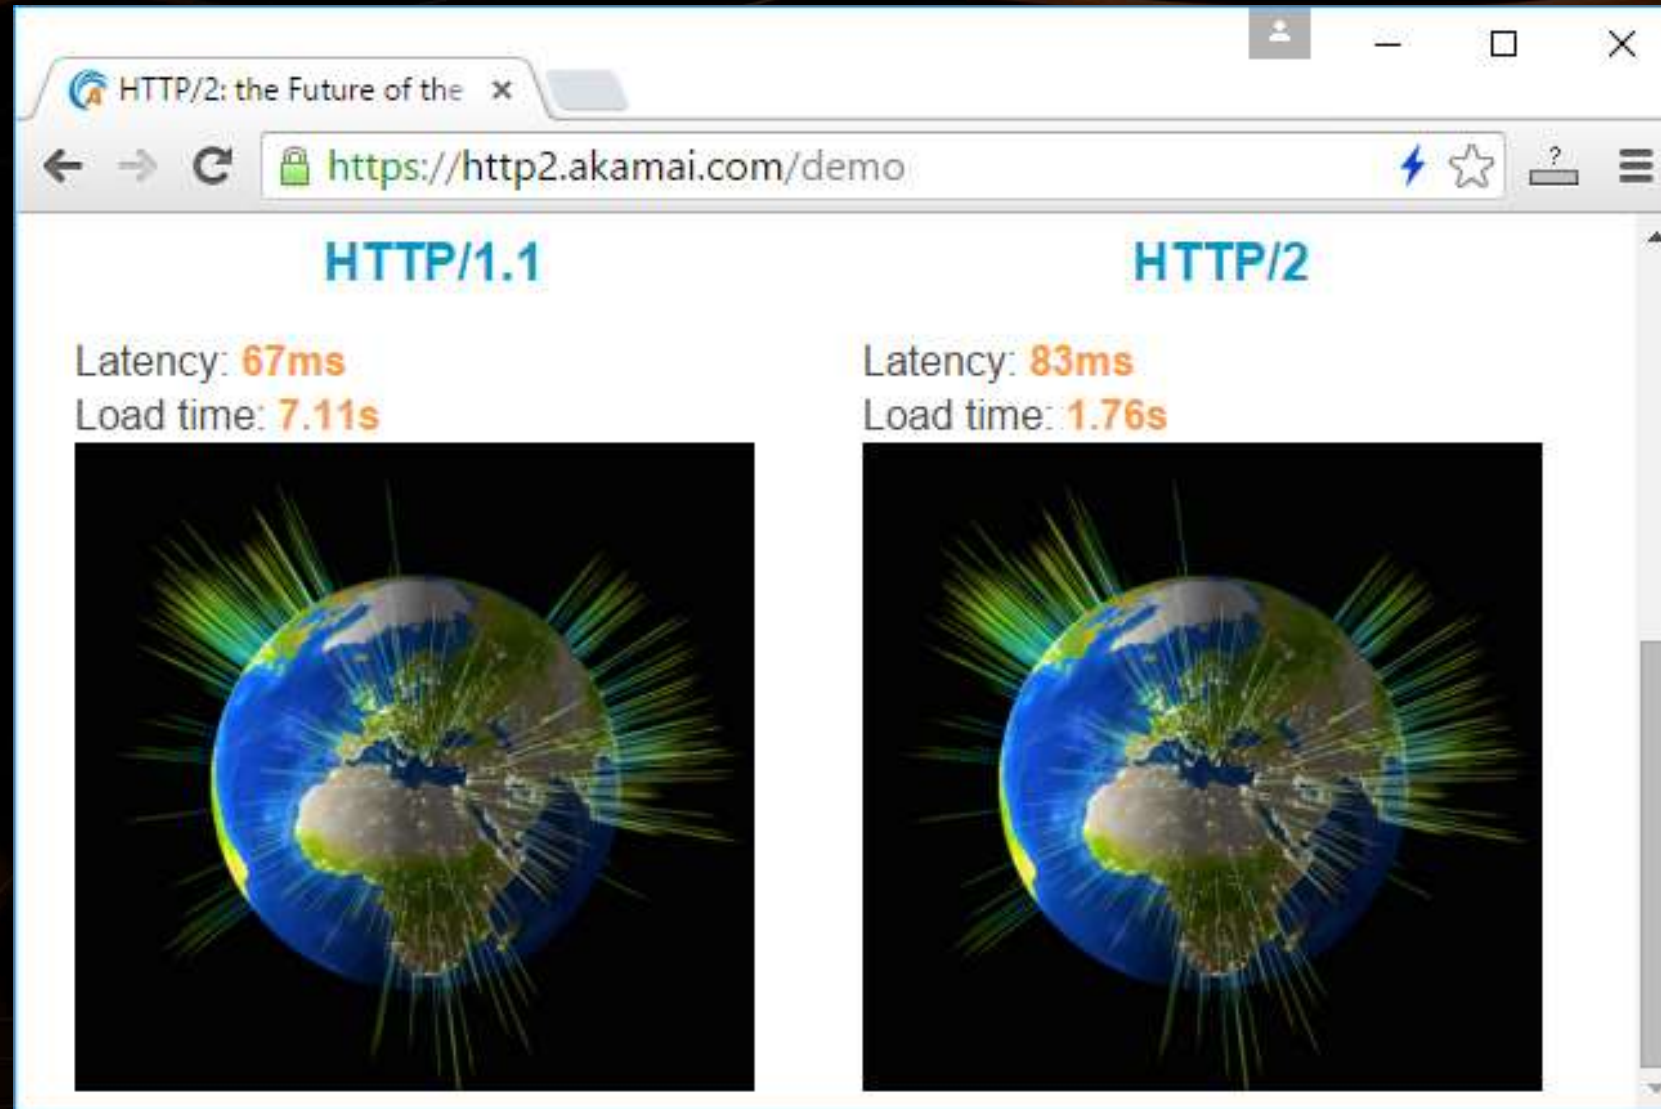
```
t=51441 [st= 0] +HTTP2_SESSION [dt=?]
--> host = "localhost:8443"
--> proxy = "DIRECT"
t=51441 [st= 0] HTTP2_SESSION_INITIALIZED
--> protocol = "h2"
--> source_dependency = 58869
```

# H2I

- **H2I** is an interactive HTTP/2 ("h2") console debugger
  - Something like **telnet** to HTTP/2 server
- <https://github.com/bradfitz/http2/tree/master/h2i>
  - Written in GO
- Send / receive raw HTTP/2 frames
  - Send **PING, SETTINGS, HEADERS** frames
  - Receive any type of frame



# Akamai HTTP/2 Performance Test





# HTTP/2 from Dev Perspective

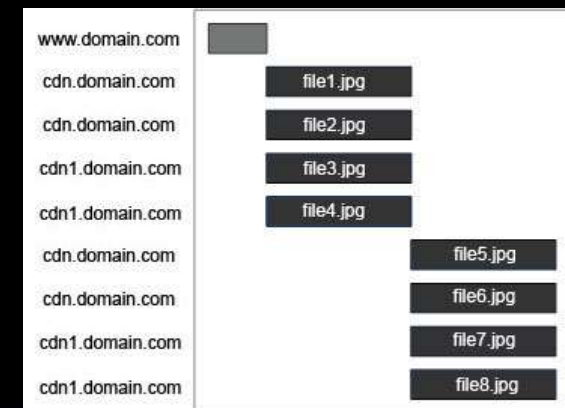
## How It Changes Developer's Life?

# Say "Goodbye" To Many Hacks

- No more image sprites
  - Use separate images
- No more CSS / JavaScript combining
  - Use separate CSS / JS files + server push
- No more domain sharding
  - Still might boost performance by distributing the server-side load
- HTTP APIs can be finer-grained without sacrificing performance



YUI Compressor





# Lots of Tweaking for Developers / Admins

- Prioritization
  - Developers may specify resource prioritization for faster page load
  - Web servers / browsers may automatically set priorities
- When to **push**?
  - Developers may **push resources** for faster page load
  - Web servers may push resources transparently of developers
- How to optimize the **flow control** ?
  - Developers may change the flow-control of data streams



# *jetty://*

## Jetty and HTTP/2

### HTTP/2 for Java Developers

# Jetty and HTTP/s


- **Jetty** is a Java HTTP (Web) server and Servlet engine
  - Developed under the Eclipse open-source project
  - Supports HTTP/2, WebSocket, OSGi, JMX, JNDI, JAAS
- Enable HTTP/2 in Jetty:

```
$ java -jar $JETTY_HOME/start.jar --add-to-startd=http2
```

```
$ java -jar $JETTY_HOME/start.jar
...
2015-06-17 14:16:12.549:INFO:oejs.ServerConnector:main: Started
ServerConnector@34c9c77f{HTTP/1.1,[http/1.1]}{0.0.0.0:8080}
2015-06-17 14:16:12.782:INFO:oejs.ServerConnector:main: Started
ServerConnector@711f39f9{SSL,[ssl, alpn, h2, h2-17, http/1.1]}{0.0.0.0:8443}
...
```

# Jetty Smart Push

- Jetty allows configuring a server push strategy
- **ReferrerPushStrategy** auto-learns dependencies from referrer



## 1. Server observes incoming traffic

- Build a dependency model based on Referer*
- e.g. `index.html` → `{style.css, app.js}`*

## 2. Server initiates push for learned dependencies

- new client → GET `index.html`*
- server → Push `style.css`, `app.js`*

*Lots of room for experimentation + innovation!*

# Jetty HTTP/2 Client

- Jetty provides **http2-client**
  - Implementation of HTTP/2 client with a low level HTTP/2 API, dealing with HTTP/2 streams, frames, etc.

```
SslContextFactory sslContextFactory = new SslContextFactory();
HttpClient httpClient = new HttpClient(sslContextFactory);
httpClient.start();
ContentResponse response =
    httpClient.GET("https://http2.akamai.com");
...
```

- Learn more at <http://www.eclipse.org/jetty/documentation/current/http-client-api.html>



# Questions?



# License

- This course (slides, examples, demos, videos, homework, etc.) is licensed under the "Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International" license



# Free Trainings @ Software University

- Software University Foundation – [softuni.org](https://softuni.org)
- Software University – High-Quality Education, Profession and Job for Software Developers
  - [softuni.bg](https://softuni.bg)
- Software University @ Facebook
  - [facebook.com/SoftwareUniversity](https://facebook.com/SoftwareUniversity)
- Software University @ YouTube
  - [youtube.com/SoftwareUniversity](https://youtube.com/SoftwareUniversity)
- Software University Forums – [forum.softuni.bg](https://forum.softuni.bg)

